

Seatwork 10.1 Case Study: Improving RT-IoT2022 Analysis

Preprocessing Dataset

Extract

```
In [192... # extract dataset and import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
IoT = pd.read_csv('datasets/RT_IOT2022.csv')
```

```
In [193... # check for duplicate values
IoT.duplicated()
```

```
Out[193... 0      False
1      False
2      False
3      False
4      False
...
123112  False
123113  False
123114  False
123115  False
123116  False
Length: 123117, dtype: bool
```

```
In [194... # check for missing values
IoT.isna().sum()
```

```
Out[194... no                0
id.orig_p            0
id.resp_p            0
proto                0
service              0
..
idle.std             0
fwd_init_window_size 0
bwd_init_window_size 0
fwd_last_window_size 0
Attack_type          0
Length: 85, dtype: int64
```

Transform

```
In [195... # I'll filter the dataset so that I can only see the columns I'll need for MY an
IoT.columns
```

```

Out[195... Index(['no', 'id.orig_p', 'id.resp_p', 'proto', 'service', 'flow_duration',
                'fwd_pkts_tot', 'bwd_pkts_tot', 'fwd_data_pkts_tot',
                'bwd_data_pkts_tot', 'fwd_pkts_per_sec', 'bwd_pkts_per_sec',
                'flow_pkts_per_sec', 'down_up_ratio', 'fwd_header_size_tot',
                'fwd_header_size_min', 'fwd_header_size_max', 'bwd_header_size_tot',
                'bwd_header_size_min', 'bwd_header_size_max', 'flow_FIN_flag_count',
                'flow_SYN_flag_count', 'flow_RST_flag_count', 'fwd_PSH_flag_count',
                'bwd_PSH_flag_count', 'flow_ACK_flag_count', 'fwd_URG_flag_count',
                'bwd_URG_flag_count', 'flow_CWR_flag_count', 'flow_ECE_flag_count',
                'fwd_pkts_payload.min', 'fwd_pkts_payload.max', 'fwd_pkts_payload.tot',
                'fwd_pkts_payload.avg', 'fwd_pkts_payload.std', 'bwd_pkts_payload.min',
                'bwd_pkts_payload.max', 'bwd_pkts_payload.tot', 'bwd_pkts_payload.avg',
                'bwd_pkts_payload.std', 'flow_pkts_payload.min',
                'flow_pkts_payload.max', 'flow_pkts_payload.tot',
                'flow_pkts_payload.avg', 'flow_pkts_payload.std', 'fwd_iat.min',
                'fwd_iat.max', 'fwd_iat.tot', 'fwd_iat.avg', 'fwd_iat.std',
                'bwd_iat.min', 'bwd_iat.max', 'bwd_iat.tot', 'bwd_iat.avg',
                'bwd_iat.std', 'flow_iat.min', 'flow_iat.max', 'flow_iat.tot',
                'flow_iat.avg', 'flow_iat.std', 'payload_bytes_per_second',
                'fwd_subflow_pkts', 'bwd_subflow_pkts', 'fwd_subflow_bytes',
                'bwd_subflow_bytes', 'fwd_bulk_bytes', 'bwd_bulk_bytes',
                'fwd_bulk_packets', 'bwd_bulk_packets', 'fwd_bulk_rate',
                'bwd_bulk_rate', 'active.min', 'active.max', 'active.tot', 'active.avg',
                'active.std', 'idle.min', 'idle.max', 'idle.tot', 'idle.avg',
                'idle.std', 'fwd_init_window_size', 'bwd_init_window_size',
                'fwd_last_window_size', 'Attack_type'],
                dtype='object')

```

```

In [196... IOT_filtered = IoT[
    ['no',
     'id.orig_p',
     'id.resp_p',
     'proto',
     'service',
     'flow_duration',
     'Attack_type',
     'fwd_pkts_per_sec',
     'bwd_pkts_per_sec',
     'payload_bytes_per_second',
     'fwd_iat.avg',
     'bwd_iat.avg',
     'flow_SYN_flag_count',
     'flow_RST_flag_count',
     'fwd_PSH_flag_count'
    ]
]
IOT_filtered.dtypes

```

```
Out[196... no int64
id.orig_p int64
id.resp_p int64
proto object
service object
flow_duration float64
Attack_type object
fwd_pkts_per_sec float64
bwd_pkts_per_sec float64
payload_bytes_per_second float64
fwd_iat.avg float64
bwd_iat.avg float64
flow_SYN_flag_count int64
flow_RST_flag_count int64
fwd_PSH_flag_count int64
dtype: object
```

Change data types for those with dtype = object.

'proto', 'service', and 'Attack_type' are categorical

I can use a dictionary to change them all at once

[Reference](#)

```
In [197... IOT_filtered_1 = IOT_filtered.astype({
    'proto': 'category',
    'service': 'category',
    'Attack_type': 'category'
})
IOT_filtered_1.dtypes
```

```
Out[197... no int64
id.orig_p int64
id.resp_p int64
proto category
service category
flow_duration float64
Attack_type category
fwd_pkts_per_sec float64
bwd_pkts_per_sec float64
payload_bytes_per_second float64
fwd_iat.avg float64
bwd_iat.avg float64
flow_SYN_flag_count int64
flow_RST_flag_count int64
fwd_PSH_flag_count int64
dtype: object
```

Load

I'll analyze the dataset by each Attack_type to delve into trends within each type.

But first, I need to group each entry based on Attack_type,
so for 1 Attack_type, there would be 1 dataframe created.

Identify what the Attack types are in the first place.

I can also count how many entries per category. Two birds, one stone.

[Reference](#)

```
In [198... IOT_filtered_1.value_counts(subset = 'Attack_type', normalize = False)
```

```
Out[198... Attack_type
DOS_SYN_Hping          94659
Thing_Speak            8108
ARP_poisoning          7750
MQTT_Publish           4146
NMAP_UDP_SCAN          2590
NMAP_XMAS_TREE_SCAN    2010
NMAP_OS_DETECTION      2000
NMAP_TCP_scan          1002
DDOS_Slowloris         534
Wipro_bulb             253
Metasploit_Brute_Force_SSH 37
NMAP_FIN_SCAN          28
Name: count, dtype: int64
```

Now, we can make dataframes for each Attack_type category

I discovered that aside from using df.query, you can also use [this](#):
df.query seems easier to remember though

```
In [259... #for DOS_SYN_Hping
ATK1 = IOT_filtered_1.loc[IOT_filtered_1['Attack_type'] == 'DOS_SYN_Hping']
```

```
In [200... #for DDOS_Slowloris
ATK9 = IOT_filtered_1.query(' Attack_type == "DDOS_Slowloris" ')
```

```
In [201... #for Wipro_bulb
ATK10 = IOT_filtered_1.query(' Attack_type == "Wipro_bulb" ')
```

```
In [202... #for Metasploit_Brute_Force_SSH
ATK11 = IOT_filtered_1.query(' Attack_type == "Metasploit_Brute_Force_SSH" ')
```

```
In [203... #for NMAP_FIN_SCAN
ATK12 = IOT_filtered_1.query(' Attack_type == "NMAP_FIN_SCAN" ')
```

What is the distribution of the Attack_type classes (normal vs. various attacks), and what percentage of the 123,117 instances does each class comprise?

References

[value_counts for a column](#)

[value_counts for a dataframe](#)

[making a new dataframe](#)

```
In [229... # the normalize parameter dictates whether proportions or frequencies will be re
Frequency = IOT_filtered_1.value_counts(subset = 'Attack_type', normalize = False)
Percentage = IOT_filtered_1["Attack_type"].value_counts(normalize=True) * 100 #s

# create a new dataframe, showing the Frequency and Percentage as columns
dist = pd.DataFrame({
    "Frequency": Frequency,
    "Percentage(%)": Percentage
```

```

    }
)

#Reset index so 'Attack_type' becomes a column instead of the index
dist = dist.reset_index()
dist.columns = ['Attack_type', 'Frequency', 'Percentage(%)']
dist.head(12)

```

Out[229...

	Attack_type	Frequency	Percentage(%)
0	DOS_SYN_Hping	94659	76.885402
1	Thing_Speak	8108	6.585606
2	ARP_poisoning	7750	6.294825
3	MQTT_Publish	4146	3.367528
4	NMAP_UDP_SCAN	2590	2.103690
5	NMAP_XMAS_TREE_SCAN	2010	1.632593
6	NMAP_OS_DETECTION	2000	1.624471
7	NMAP_TCP_scan	1002	0.813860
8	DDOS_Slowloris	534	0.433734
9	Wipro_bulb	253	0.205496
10	Metasploit_Brute_Force_SSH	37	0.030053
11	NMAP_FIN_SCAN	28	0.022743

In [239...

```

# Get the Last 4 rows
least_common = dist.tail(4)

# Step 2: Calculate total percentage
total = dist['Frequency'].sum()
least_common_percentage_sum = (least_common['Frequency'].sum() / total) * 100

# Step 3: Print the result
print(f"Total percentage of the 4 least common attack types: {least_common_perce

```

Total percentage of the 4 least common attack types: 0.69%

We can represent this with a pie chart

In [242...

```

import matplotlib.pyplot as plt

# Sort and get the top 3 frequencies
top3_indices = dist['Frequency'].nlargest(3).index.tolist()

# Convert frequencies to List
frequencies = dist['Frequency'].tolist()

# Compute percentages
total = sum(frequencies)
percentages = [(f / total) * 100 for f in frequencies]

# Create autopct Labels: only top 3 get Labels
autopct_labels = [f"{p:.1f}%" if i in top3_indices else "" for i, p in enumerate

```

```

# Plot
fig, ax = plt.subplots(figsize=(5, 5))
wedges, texts, autotexts = ax.pie(
    frequencies,
    startangle=90,
    labels=None,
    autopct=lambda pct: autopct_labels.pop(0) # Use precomputed List
)


# Add Legend
ax.legend(dist['Attack_type'],
          title="Attack Types",
          loc="center left",
          bbox_to_anchor=(1, 0.5),
          fontsize=9)

# Title
ax.set_title('Attack Type Distribution')

# Show
plt.show()

```

Attack Type Distribution



How do the categorical features proto (protocol) and service vary across different attack types and normal traffic patterns?

"service"

In [254... ATK9['service'].value_counts()

```
Out[254... service
http      523
-          6
dns        3
dhcp       2
irc        0
mqtt       0
ntp        0
radius     0
ssh        0
ssl        0
Name: count, dtype: int64
```

```
In [255... ATK10['service'].value_counts()
```

```
Out[255... service
ssl        107
dns         53
-           45
irc         43
dhcp         5
http         0
mqtt         0
ntp         0
radius       0
ssh          0
Name: count, dtype: int64
```

```
In [256... ATK11['service'].value_counts()
```

```
Out[256... service
ssh         28
dns          8
http         1
-            0
dhcp         0
irc          0
mqtt         0
ntp          0
radius       0
ssl          0
Name: count, dtype: int64
```

```
In [257... ATK12['service'].value_counts()
```

```
Out[257... service
-          24
dns         3
http        1
dhcp        0
irc         0
mqtt        0
ntp         0
radius      0
ssh         0
ssl         0
Name: count, dtype: int64
```

```
In [261... # Get value_counts as percentages
atk9_perc = ATK9['service'].value_counts(normalize=True).rename('DDOS_Slowloris')
```

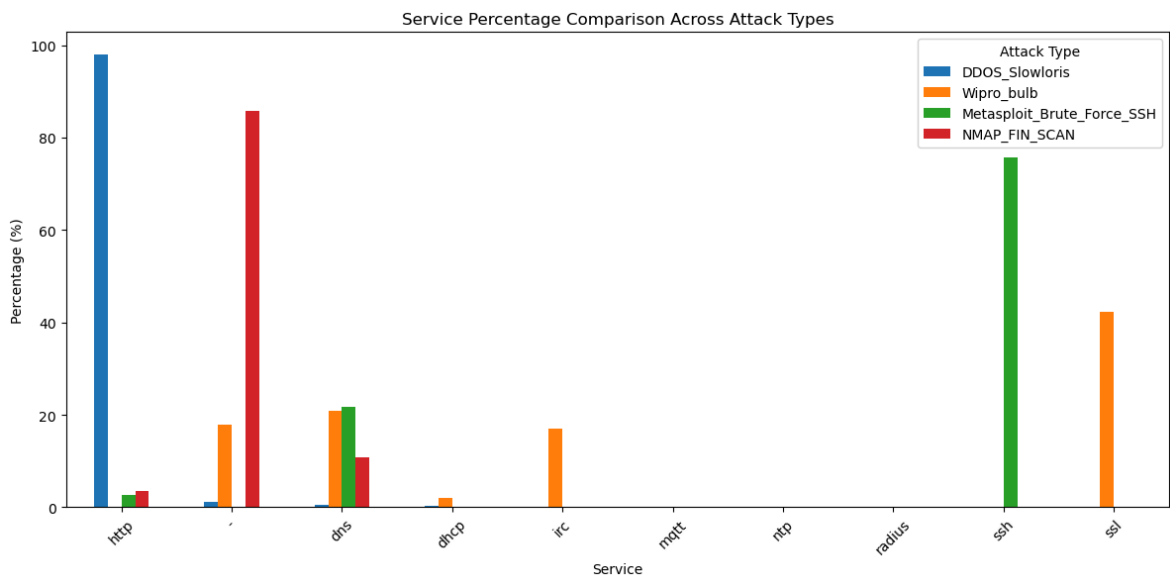
```

atk10_perc = ATK10['service'].value_counts(normalize=True).rename('Wipro_bulb')
atk11_perc = ATK11['service'].value_counts(normalize=True).rename('Metasploit_Br')
atk12_perc = ATK12['service'].value_counts(normalize=True).rename('NMAP_FIN_SCAN')

# Combine into one DataFrame
combined_percentages = pd.concat([atk9_perc, atk10_perc, atk11_perc, atk12_perc])

# Plot
combined_percentages.plot(kind='bar', figsize=(12, 6))
plt.title('Service Percentage Comparison Across Attack Types')
plt.xlabel('Service')
plt.ylabel('Percentage (%)')
plt.legend(title='Attack Type')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



"proto"

In [249...] `ATK9['proto'].value_counts()`

Out[249...] proto
tcp 529
udp 5
icmp 0
Name: count, dtype: int64

In [250...] `ATK10['proto'].value_counts()`

Out[250...] proto
tcp 182
udp 69
icmp 2
Name: count, dtype: int64

In [251...] `ATK11['proto'].value_counts()`


```
Out[251...] proto
tcp      29
udp       8
icmp      0
Name: count, dtype: int64
```

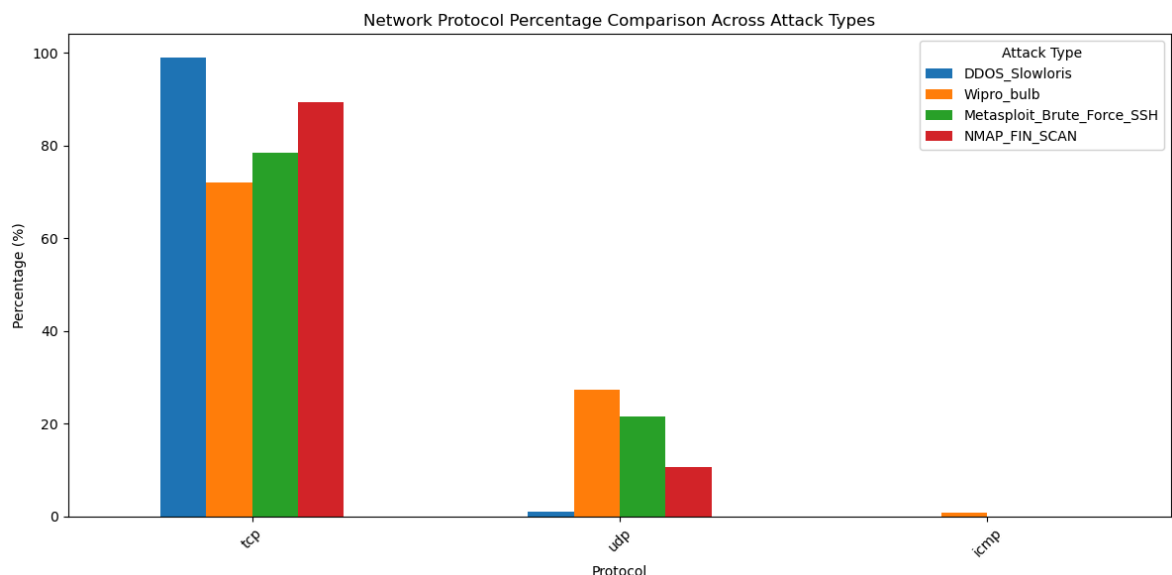
```
In [252...] ATK12['proto'].value_counts()
```

```
Out[252...] proto
tcp      25
udp       3
icmp      0
Name: count, dtype: int64
```

```
In [262...] # Get value_counts as percentages
atk9_perc = ATK9['proto'].value_counts(normalize=True).rename('DDOS_Slowloris')
atk10_perc = ATK10['proto'].value_counts(normalize=True).rename('Wipro_bulb') *
atk11_perc = ATK11['proto'].value_counts(normalize=True).rename('Metasploit_Brut')
atk12_perc = ATK12['proto'].value_counts(normalize=True).rename('NMAP_FIN_SCAN')

# Combine into one DataFrame
combined_percentages = pd.concat([atk9_perc, atk10_perc, atk11_perc, atk12_perc])

# Plot
combined_percentages.plot(kind='bar', figsize=(12, 6))
plt.title('Network Protocol Percentage Comparison Across Attack Types')
plt.xlabel('Protocol')
plt.ylabel('Percentage (%)')
plt.legend(title='Attack Type')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



What are the mean and standard deviation of flow_duration for each Attack_type, and are differences statistically significant?

```
In [206...] # run .mean and .stdev
```

I'll only analyze the Attack_types that occurred less than 1000 times.

for DDOS_Slowloris

```
In [207... print("DDOS_Slowloris")
mean = ATK9['flow_duration'].mean()
median = ATK9['flow_duration'].median()
stdev = ATK9['flow_duration'].std()

Skewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

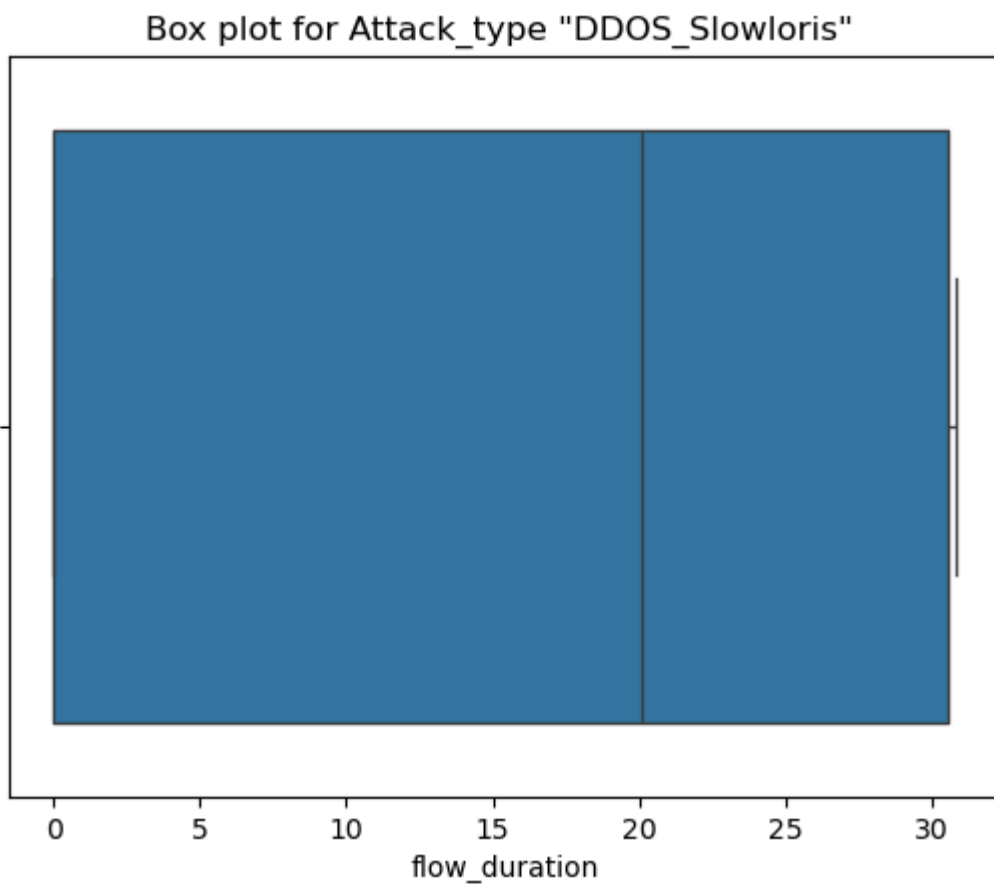
ATK9['flow_duration'].describe()
```

DDOS_Slowloris

Skewness = -1.1447559217527645

```
Out[207... count    534.000000
mean      14.699148
std       14.124797
min        0.000000
25%        0.003817
50%       20.088963
75%       30.557698
max       30.870463
Name: flow_duration, dtype: float64
```

```
In [208... sns.boxplot(data=ATK9, x = 'flow_duration')
plt.title('Box plot for Attack_type "DDOS_Slowloris" ')
plt.show()
```



Box plot shows no outliers.

No need to filter the dataframe.

for Wipro_bulb

```
In [210... print("Wipro_bulb")
mean = ATK10['flow_duration'].mean()
median = ATK10['flow_duration'].median()
stdev = ATK10['flow_duration'].std()

kewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

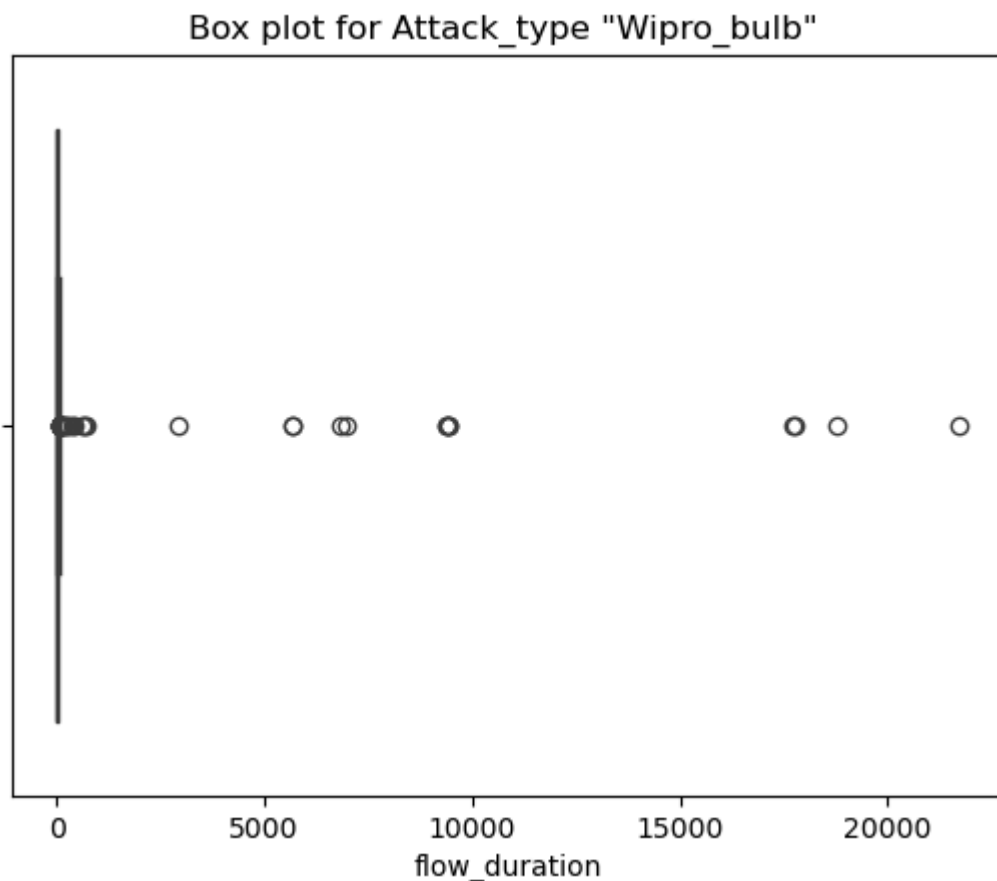
ATK10['flow_duration'].describe()
```

Wipro_bulb

Skewness = -1.1447559217527645

```
Out[210... count    253.000000
mean      586.845727
std       2738.891637
min        0.000000
25%        0.027666
50%        0.803326
75%       31.285402
max      21728.335578
Name: flow_duration, dtype: float64
```

```
In [211... sns.boxplot(data=ATK10, x = 'flow_duration')
plt.title('Box plot for Attack_type "Wipro_bulb" ')
plt.show()
```



The dataframe is slightly skewed positively

The box plot also shows that a number of outliers affect the mean and consequently, the standard deviation.

```
In [212... # Flagging outliers using IQR
# I'll use this method to weed out the outliers since the data points are slight
Q1 = ATK10['flow_duration'].quantile(0.25)
Q3 = ATK10['flow_duration'].quantile(0.75)
IQR = Q3 - Q1

# Identify outliers
""" If the value is
    less than the first quartile minus 1.5 times the Interquartile range or
    greater than the 3rd quartile plus 1.5 times the Interquartile range,
    they're considered an outlier
    References I've read used 1.5, so I also used it. ^_^
"""
outliers = (ATK10['flow_duration'] < (Q1 - 1.5* IQR)) | (ATK10['flow_duration']

# Check the flagged outliers
outliers.value_counts()
```

```
Out[212... flow_duration
False      221
True        32
Name: count, dtype: int64
```

Let's make a dataframe that has less of these outliers

```
In [213... ATK10_no_outliers = ATK10[~outliers] # Negate the outliers to keep non-outliers
print("Wipro_bulb | less outliers")
mean = ATK10_no_outliers['flow_duration'].mean()
median = ATK10_no_outliers['flow_duration'].median()
stdev = ATK10_no_outliers['flow_duration'].std()

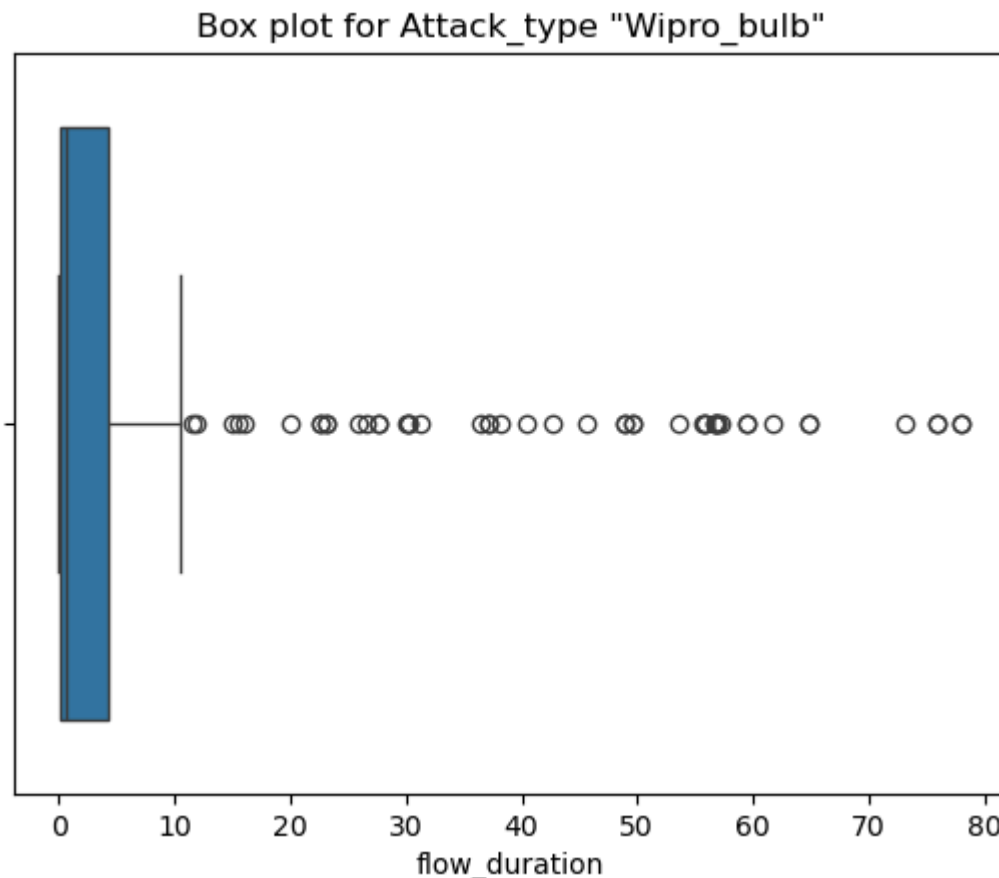
Skewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

ATK10_no_outliers['flow_duration'].describe()
```

```
Wipro_bulb | less outliers
Skewness = 1.459505897397956
```

```
Out[213... count      221.000000
mean       10.370136
std        19.959788
min         0.000000
25%         0.021565
50%         0.659660
75%         4.182818
max        78.044458
Name: flow_duration, dtype: float64
```

```
In [214... sns.boxplot(data=ATK10_no_outliers, x = 'flow_duration')
plt.title('Box plot for Attack_type "Wipro_bulb"')
plt.show()
```



for Metasploit_Brute_Force_SSH

```
In [215...] print("Metasploit_Brute_Force_SSH ")
mean = ATK11['flow_duration'].mean()
median = ATK11['flow_duration'].median()
stdev = ATK11['flow_duration'].std()

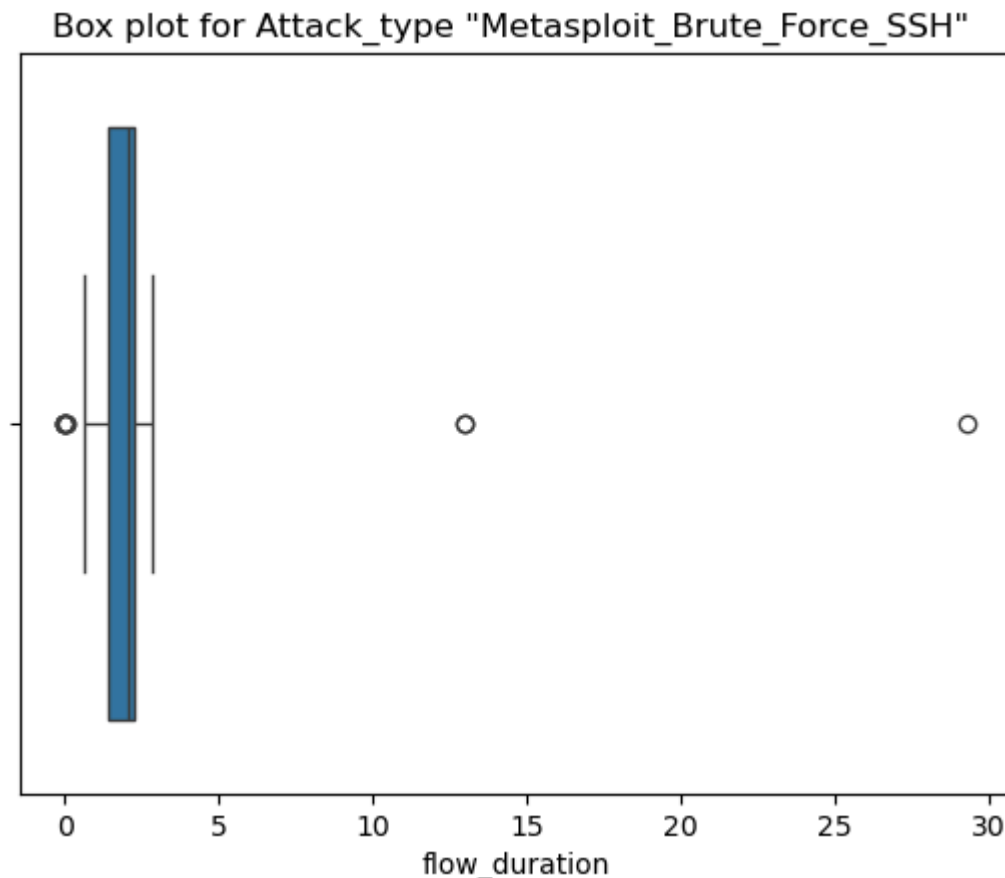
kewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

ATK11['flow_duration'].describe()
```

```
Metasploit_Brute_Force_SSH
Skewness = 1.459505897397956
```

```
Out[215...] count    37.000000
mean      3.006557
std       5.210286
min       0.000000
25%      1.417588
50%      2.030317
75%      2.258765
max      29.289262
Name: flow_duration, dtype: float64
```

```
In [216...] sns.boxplot(data=ATK11, x = 'flow_duration')
plt.title('Box plot for Attack_type "Metasploit_Brute_Force_SSH" ')
plt.show()
```



The dataframe is skewed positively

The box plot also shows that a number of outliers affect the mean and consequently, the standard deviation.

```
In [217... # Flagging outliers using IQR
# I'll use this method to weed out the outliers since the data points are slight
Q1 = ATK11['flow_duration'].quantile(0.25)
Q3 = ATK11['flow_duration'].quantile(0.75)
IQR = Q3 - Q1

# Identify outliers
""" If the value is
    less than the first quartile minus 1.5 times the Interquartile range or
    greater than the 3rd quartile plus 1.5 times the Interquartile range,
    they're considered an outlier
    References I've read used 1.5, so I also used it. ^_^
"""
outliers = (ATK11['flow_duration'] < (Q1 - 1.5* IQR)) | (ATK11['flow_duration']
# Check the flagged outliers
outliers.value_counts()
```

```
Out[217... flow_duration
False      28
True        9
Name: count, dtype: int64
```

```
In [218... ATK11_no_outliers = ATK11[~outliers] # Negate the outliers to keep non-outliers
print("Wipro_bulb | less outliers")
mean = ATK11_no_outliers['flow_duration'].mean()
median = ATK11_no_outliers['flow_duration'].median()
```

```

stdev = ATK11_no_outliers['flow_duration'].std()

Skewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

ATK11_no_outliers['flow_duration'].describe()

```

Wipro_bulb | less outliers

Skewness = -0.3868506068156085

```

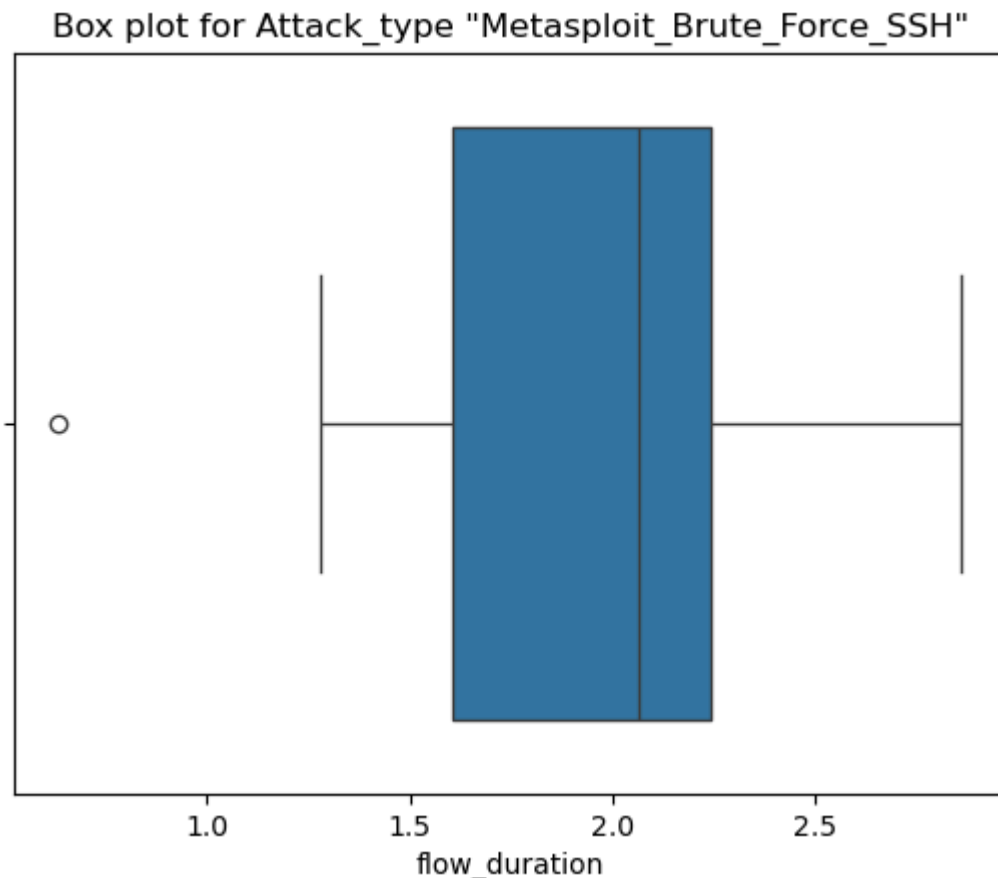
Out[218...] count    28.000000
            mean     1.999530
            std      0.508938
            min      0.633988
            25%      1.607395
            50%      2.065157
            75%      2.242271
            max      2.861502
            Name: flow_duration, dtype: float64

```

```

In [219...] sns.boxplot(data=ATK11_no_outliers, x = 'flow_duration')
plt.title('Box plot for Attack_type "Metasploit_Brute_Force_SSH"')
plt.show()

```



for NMAP_FIN_SCAN

```

In [220...] print("NMAP_FIN_SCAN")
mean = ATK12['flow_duration'].mean()
median = ATK12['flow_duration'].median()
stdev = ATK12['flow_duration'].std()

kewness = 3*(mean - median) / stdev
print("Skewness = ", Skewness)

```

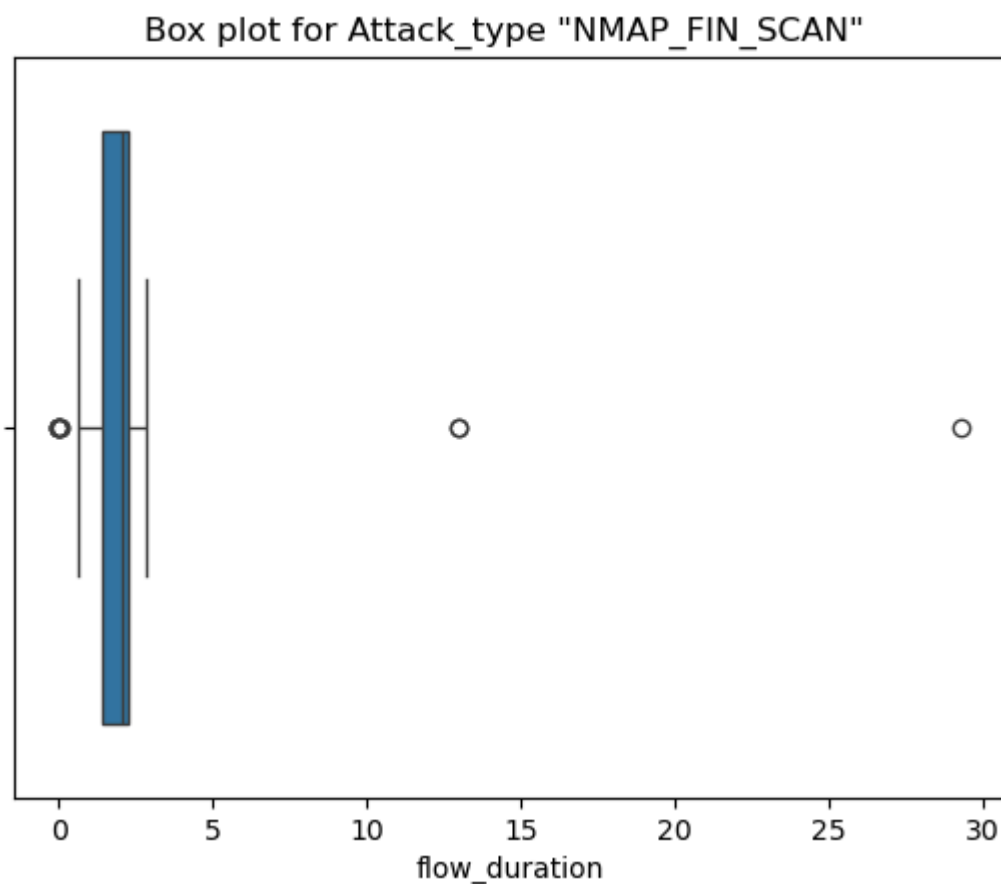
```
ATK12['flow_duration'].describe()
```

NMAP_FIN_SCAN

Skewness = -0.3868506068156085

```
Out[220...] count    28.000000
            mean     0.023614
            std      0.108791
            min      0.000000
            25%      0.000000
            50%      0.000000
            75%      0.000000
            max      0.575884
            Name: flow_duration, dtype: float64
```

```
In [221...] sns.boxplot(data=ATK11, x = 'flow_duration')
            plt.title('Box plot for Attack_type "NMAP_FIN_SCAN" ')
            plt.show()
```



The dataframe is slightly skewed negatively

The box plot also shows that a number of outliers affect the mean and consequently, the standard deviation. However, the quartiles are 0, so filtering by using the IQR is not possible.

Now that the mean and standard deviation for the selected attack types have been determined, I think we can graph them now to visualize their differences.

```
In [264...] #plotting the mean amongst the 4 attack types
            import matplotlib.pyplot as plt

            # Calculate means for flow_duration in each dataframe
```



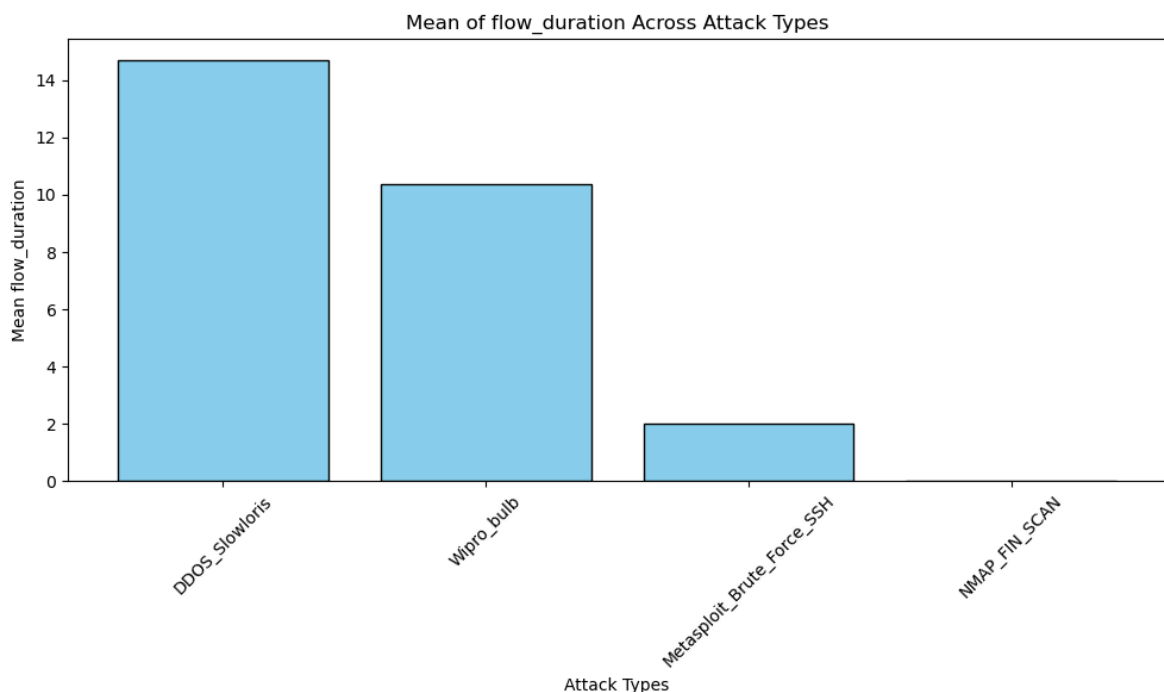
```

means = [
    ATK9['flow_duration'].mean(),
    ATK10_no_outliers['flow_duration'].mean(),
    ATK11_no_outliers['flow_duration'].mean(),
    ATK12['flow_duration'].mean()
]

# Names of the attack types
attack_types = ['DDOS_Slowloris', 'Wipro_bulb', 'Metasploit_Brute_Force_SSH', 'N

# Create the plot for means
fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(attack_types, means, color='skyblue', edgecolor='black')
ax.set_title('Mean of flow_duration Across Attack Types')
ax.set_xlabel('Attack Types')
ax.set_ylabel('Mean flow_duration')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



In [265...

```

#plotting the stdev amongst the 4 attack types
import matplotlib.pyplot as plt

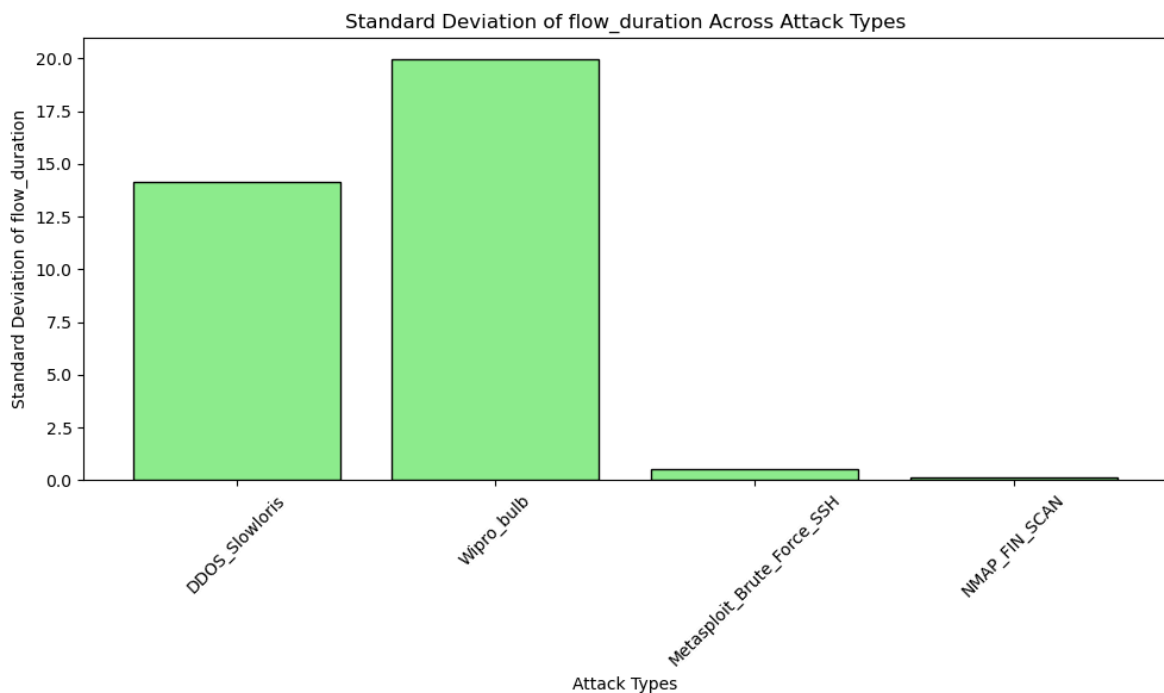
# Calculate standard deviations for flow_duration in each dataframe
stdevs = [
    ATK9['flow_duration'].std(),
    ATK10_no_outliers['flow_duration'].std(),
    ATK11_no_outliers['flow_duration'].std(),
    ATK12['flow_duration'].std()
]

# Names of the attack types
attack_types = ['DDOS_Slowloris', 'Wipro_bulb', 'Metasploit_Brute_Force_SSH', 'N

# Create the plot for standard deviations
fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(attack_types, stdevs, color='lightgreen', edgecolor='black')
ax.set_title('Standard Deviation of flow_duration Across Attack Types')

```

```
ax.set_xlabel('Attack Types')
ax.set_ylabel('Standard Deviation of flow_duration')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Which continuous features (e.g., fwd_pkts_per_sec, bwd_pkts_per_sec, payload_bytes_per_second) exhibit the highest correlation with specific attack classes?

In []:

How do time-based features like fwd_iat.avg and bwd_iat.avg (mean inter-arrival times) differ between various attack types and normal traffic?

In []:

Which network flag counts (e.g., flow_SYN_flag_count, flow_RST_flag_count, fwd_PSH_flag_count) are most indicative of specific intrusion patterns?

In []:

In []: